

Marc Chengliang Zhang

PRESENT POSITION

Ph.D Candidate
Dept. of Computer Science & Engineering
Hong Kong University of Science and Technology
Clear Water Bay, Hong Kong

Office: Room CYT 3007
Phone: +852-6216-2287
Email: czhangbn@cse.ust.hk
Web: <https://marcoszh.github.io/>

RESEARCH INTERESTS

My interests cover **big data analytics systems** and **cloud computing**, with a special focus on **machine learning systems**. I enjoy identifying fundamental system design and performance issues in large-scale ML systems for both training and inference, and searching for general and efficient solutions.

EDUCATION

Hong Kong University of Science and Technology, Hong Kong SAR
Department of Computer Science and Engineering

- ◇ **Ph.D.** Computer Science and Engineering September 2016 - present
 - ◇ Supervisor: [Wei Wang](#)
 - ◇ [HK PhD Fellowship](#) recipient: prestigious and highly selective fellowship

Harbin Institute of Technology, Harbin, China
School of Computer Science and Technology

- ◇ **B.Eng.** Software Engineering September 2012 - June 2016
 - ◇ Honors: National Scholarship (Top 2%), People's Scholarship, Fuji Xerox Scholarship

PUBLICATIONS

[Chengliang Zhang](#), Minchen Yu, Wei Wang, Feng Yan, "[MARk: Exploiting Cloud Services for Cost-Effective, SLO-Aware Machine Learning Inference Serving](#)," in the *Proceedings of USENIX Annual Technical Conference (ATC'19)*, Renton, WA, July 2018 (20% acceptance rate).

[Chengliang Zhang](#), Huangshi Tian, Wei Wang, Feng Yan, "[Stay Fresh: Speculative Synchronization for Fast Distributed Machine Learning](#)," in the *Proceedings of IEEE International Conference on Distributed Computing Systems (ICDCS'18)*, Vienna, Austria, July 2018 (20% acceptance rate).

Preprints

[Chengliang Zhang](#), Minchen Yu, Wei Wang, Feng Yan, "Towards Cost-Effective and SLO-Aware Machine Learning Inference Serving on Public Cloud," to be submitted to *IEEE Transactions on Parallel and Distributed Systems*.

Yinghao Yu, [Chengliang Zhang](#), Wei Wang, Jun Zhang, Khaled Letaief, "Towards Dependency-Aware Cache Management for Data Analytics Applications," submitted to *IEEE Transactions on Cloud Computing*, currently under review.

RESEARCH EXPERIENCE

Fast Secure Federated Learning System

- Inter-enterprise federated learning with Homomorphic Encryption
- Accelerate training and inference of Secure Federated Learning
- Mitigate encryption and communication overhead

MArk: ML Serving on Public Cloud

- Serve machine learning inference on public cloud
- Cost-effective and SLO-aware
- Characterization of ML serving and its performance cloud services
- Combine FaaS and IaaS to reduce over-provisioning
- Characterization of hardware accelerators like GPU and TPU

Speculative Synchronization

- Distributed data parallel training
- Relaxed consistency can increase throughput but hurt update quality
- Re-synchronize if the parameter copy is too stale to produce beneficial updates

SKILLS

- | | | |
|--------------|---------------|---------|
| ◇ Python | ◇ Java, Scala | ◇ Keras |
| ◇ TensorFlow | ◇ MXNet | ◇ Spark |
| ◇ Hadoop | | |

REFERENCES

Wei Wang, Assistant Professor
 Hong Kong University of Science and Technology
 Room 3524, CSE Department, HKUST
 Clear Water Bay, Kowloon, Hong Kong
Phone: +852-2358-6972
Email: weiwa@cse.ust.hk
Web: <https://www.cse.ust.hk/weiwa/>

Feng Yan, Assistant Professor
 University of Nevada, Reno
 SEM 233, CSE Department Reno, NV 89557
Phone: +1-775-784-6448
Email: fyan@unr.edu
Web: <https://wolfweb.unr.edu/fyan/>